

Hello, Im **Gokulnath**, please find my task submission below.

🚀 FastAPI Server for Retrieval Augmented Generation (RAG) System

Task Overview

The goal was to implement a **lightweight FastAPI server** designed for a **Retrieval Augmented Generation (RAG)** system with the following requirements:

- 🗄️ **ChromaDB** for document ingestion and querying
 - 📁 Support for multiple document formats: **PDF, DOC, DOCX, TXT**
 - 🧠 Use **sentence-transformers/all-MiniLM-L6-v2** from Hugging Face for embeddings
 - ⚡ Ensure **non-blocking API endpoints** and efficient concurrency with FastAPI's async features
-

Task submission - Github repo : <https://github.com/Coding-Devil/FastAPI-Server-for-RAG/tree/main>

🔧 My Approach and Implementation

1. FastAPI Server

- Developed a FastAPI server utilizing **asynchronous capabilities** for handling non-blocking operations.
- The server exposes endpoints for:
 - 📄 **Document upload**
 - 🔍 **Querying the database**
 - 📋 **Listing available collections**

2. Document Processing

- Implemented support for various document types including **PDF, DOC, DOCX, and TXT**.
- Asynchronous file handling via **asyncio** ensures smooth, non-blocking processing.

3. Embeddings

- Chose **sentence-transformers/all-MiniLM-L6-v2** from Hugging Face for generating embeddings. This model is lightweight and performs efficiently even on **CPUs**.

4. Vector Database Integration

- Integrated **ChromaDB** as the vector database, allowing persistent storage and querying of document embeddings.







5. RAG System Implementation

- The **RAG server** retrieves relevant document chunks based on user queries, leveraging **sentence-transformers** to generate **contextual answers**.

6. Concurrency & Non-blocking Operations

- For time-consuming tasks like file uploads and document ingestion, used FastAPI's **async features** and Python's **asyncio**.
 - Background tasks handle long-running processes like document parsing and embedding generation to keep the API **responsive**.
-







Technologies Used

-  **FastAPI**: Lightweight server framework
 -  **ChromaDB**: Persistent vector storage
 -  **Sentence-Transformers**: Embedding generation
 -  **Hugging Face Inference API**: For RAG-based text generation
 -  **PyPDF2 & python-docx**: Document parsing tools
 -  **Asyncio**: For managing asynchronous operations
-

Project Structure

```
rag_server/
├── app.py           # Main FastAPI application
├── vector_db.py     # ChromaDB integration logic
├── load_data.py     # Document loading and splitting logic
├── prompts.py      # RAG prompt generation
├── utils.py        # Utility functions (e.g., async helpers)
├── ingest.py       # Document ingestion logic
├── COLLECTIONS.txt # List of document collections
└── data/           # Directory for uploaded documents
```

Meeting Expectations

-  **Lightweight Server**: FastAPI ensures efficient, lightweight performance
 -  **ChromaDB Integration**: Persistent vector storage with querying
 -  **Document Support**: Handles **PDF** and **DOCX**, extendable to other formats like **HTML**
 -  **Embeddings**: Using **sentence-transformers** for efficient embeddings
 -  **Non-blocking API**: FastAPI's async capabilities guarantee non-blocking endpoints
 -  **Concurrency**: Efficient handling of concurrent requests through **asyncio**
-

Results and Demonstration

The system successfully:

1. Uploads and processes multiple document formats.
2. Efficiently queries stored document collections using embeddings.
3. Provides accurate, context-aware responses via RAG, leveraging pre-trained NLP models.

- Is scalable and can handle various document collections with ease.

RAG Fast-API Server 1.0.0 OAS 3.1

/openapi.json

Retrieval Augmented Generation APP which lets users upload a file and get answers to questions using LLMs. This API allows you to upload documents (PDF, DOC, DOCX, TXT), query them, and get AI-generated answers based on the document content.

default

POST	/upload	Upload File	⌵
GET	/query	Query	⌵
GET	/collections	List Collections	⌵

Schemas

Body_upload_file_upload_post > Expand all object

HTTPValidationError > Expand all object

ValidationError > Expand all object

Ingest :

Retrieval Augmented Generation APP which lets users upload a file and get answers to questions using LLMs

default

GET	/	Index	⌵
-----	---	-------	---

POST	/upload	Upload File	⌵
Parameters		Cancel	Reset
Name	Description		
collection_name	test_collection		
string	(query)		
Request body required		multipart/form-data ⌵	
file * required	Choose File	Gokulnath Resume Cybermetyx.pdf	
string(\$binary)			
Servers			
These operation-level options override the global server options.			
/ ⌵			
Execute			

Upload:

Execute

Clear

Responses

Curl

```
curl -X 'POST' \
  http://127.0.0.1:8000/upload?collection_name=test_collection' \
  -H 'accept: application/json' \
  -H 'Content-Type: multipart/form-data' \
  -F 'file=@Gokulnath Resume Cybernetyx.pdf;type=application/pdf'
```

Request URL

http://127.0.0.1:8000/upload?collection_name=test_collection

Server response

Code	Details
200	<div><div>Response body</div><div><pre>{ "message": "File Gokulnath Resume Cybernetyx.pdf uploaded. Processing in the background." }</pre></div><div><div>Download</div></div></div> <div><div>Response headers</div><pre>content-length: 90 content-type: application/json date: Tue, 15 Oct 2024 20:40:51 GMT server: uvicorn</pre></div>

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

Controls Accept header.

Example Value

Schema

Query :

GET /query Query

Cancel

Parameters

Name	Description
query * required	what are his skills?
string (query)	
n_results	Number of results to return
integer (query)	2
collection_name	Name of the document collection to search
string (query)	test_collection

Servers

These operation-level options override the global server options.

/

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
  http://127.0.0.1:8000/query/query=what%20are%20his%20skills&n_results=2&collection_name=test_collection' \
  -H 'accept: application/json'
```

Output :

Responses

Curl

```
curl -X 'GET' \
  http://127.0.0.1:8000/query?query=what%20are%20his%20skills%3F&n_results=2&collection_name=test_collection' \
  -H 'accept: application/json'
```

Request URL

http://127.0.0.1:8000/query?query=what%20are%20his%20skills%3F&n_results=2&collection_name=test_collection

Server response

Code	Details
200	<div><div>Response body</div><div><pre>{ "message": "Query is what are his skills?", "llm_output": " According to the provided excerpts, Gokulnath has skills in Python, C++, web development, machine learning, and deploying scalable applications." }</pre></div><div><div>Response headers</div><div><pre>content-length: 204 content-type: application/json date: Tue, 15 Oct 2024 22:25:33 GMT server: uvicorn</pre></div></div></div>

Responses

Code	Description	Links
200	<div><div>Successful Response</div><div><div>Media type</div><div>application/json</div><div>Controls Accept header</div><div>Example Value Schema</div><div><pre>"string"</pre></div></div></div>	No links

Conclusion

In conclusion, this project demonstrates the successful creation of a robust, scalable, and efficient RAG system that adheres to all the outlined requirements. The implementation makes use of modern Python async programming, advanced NLP models from Hugging Face, and effective vector-based information retrieval via ChromaDB.